

Group B assignment 1

write a code in java for a simple word count application that counts the number of occurrences of each word in a given input set using the hadoop map reduce framework on local standalone setup

```
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {

    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable>{

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context
            ) throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
```

```
        word.set(itr.nextToken());  
        context.write(word, one);  
    }  
}  
}
```

```
public static class IntSumReducer  
    extends Reducer<Text,IntWritable,Text,IntWritable> {  
    private IntWritable result = new IntWritable();  
  
    public void reduce(Text key, Iterable<IntWritable> values,  
        Context context  
    ) throws IOException, InterruptedException {  
        int sum = 0;  
        for (IntWritable val : values) {  
            sum += val.get();  
        }  
        result.set(sum);  
        context.write(key, result);  
    }  
}
```

```
public static void main(String[] args) throws Exception {  
    Configuration conf = new Configuration();  
    Job job = Job.getInstance(conf, "word count");  
    job.setJarByClass(WordCount.class);  
    job.setMapperClass(TokenizerMapper.class);  
    job.setCombinerClass(IntSumReducer.class);  
    job.setReducerClass(IntSumReducer.class);  
}
```

```
job.setOutputKeyClass(Text.class);

job.setOutputValueClass(IntWritable.class);

FileInputFormat.addInputPath(job, new Path(args[0]));

FileOutputFormat.setOutputPath(job, new Path(args[1]));

System.exit(job.waitForCompletion(true) ? 0 : 1);

}

}
```

To run this code on a local standalone setup, you need to have Hadoop installed and properly configured on your system. Once you have Hadoop set up, you can compile the above code using `javac`, and then create a JAR file using `jar` command. After that, you can use the following command to execute the code:

```
hadoop jar WordCount.jar WordCount <input_path> <output_path>
```

Replace `<input_path>` with the path to your input file or directory, and `<output_path>` with the desired location for the output. The results will be stored in the specified output directory.

Please note that setting up Hadoop and running MapReduce jobs can be a complex process, and this example assumes you have a working Hadoop setup on your local machine. If you don't have Hadoop installed, you can try running the code in a cloud-based Hadoop environment like AWS EMR or Cloudera.